

COMS20012

Segmentation, Pages and Memory Protection

Joseph Hallett

bristol.ac.uk



What's all this about?

- Many processes run at the same time in separate address spaces
 - Isn't *virtual memory* neat?
- So how are we going to get the OS to deal with getting the right bit of memory into the right place at the right time?
- ...and what can we do with with these mechanisms once we've got them going



Here be dragons!

- This stuff is super confusing, and we're going to focus on the mechanisms for the Intel x86 architecture which is especially weird, arcane and tricky
- If in doubt, go read Volume 3, Chapters 3–5 of the *Intel 64 and IA-32 Architectures Software Developer's Manual* which gives details
- Get the broad ideas down... go into more detail if you need/want!



Segmentation

- We want different programs and tasks to run on a single processor without interfering with each other and the OS kernel...
- Code and stacks should all seem to start at the same addresses yet should really refer to different bits of memory...
- Oh and we might have more memory available than our CPU can simply address with an n -bit register...
- So how are we going to do this?



Segment Registers

- The X86 instruction set has a bunch of *segment registers*
 - As well as all the standard RAX/RBX/RIP/RBP registers...
- CS is the code segment
- DS is the data segment
- SS is the stack segment
- ES is the extra segment (*used for strings mostly*)
- FS/GS are general purpose segments
- Pointers are treated as an offset from one of these segments
- *Global Descriptor Table (GDT)* says where all these segments live
- Each segment can have a set of access permissions



And what happens with these all these segment registers?

(Basically, it all goes down a rabbit hole of tables and virtual address spaces...)

(let's focus on the left-hand side for now)

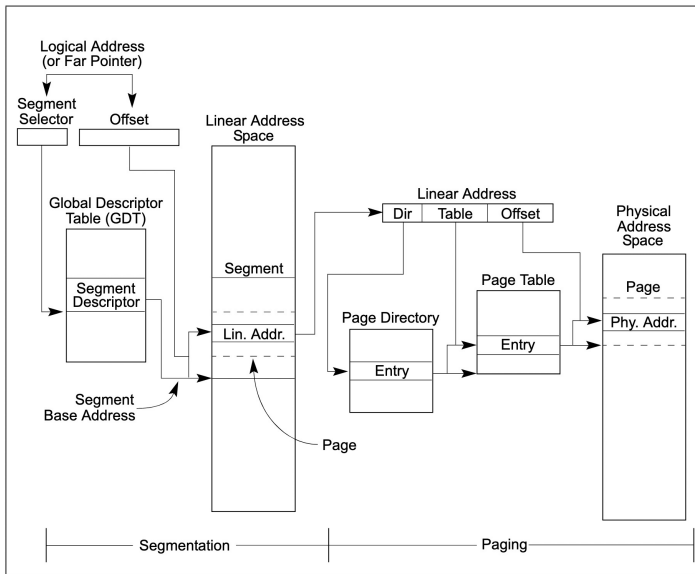


Figure 3-1. Segmentation and Paging



Flat Model

(Segments for code and data/stack... try and pretend it doesn't really exist)

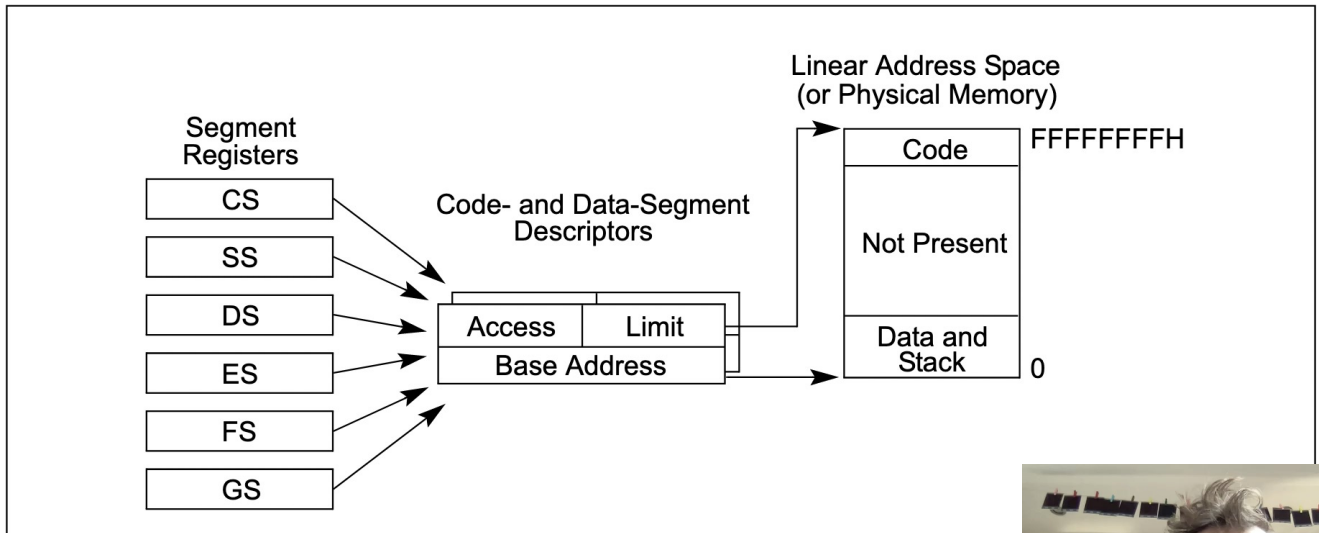


Figure 3-2. Flat Model



Protected Flat Model

(Separate segments for ring 0 and ring 3 code and data... maybe more if helpful)

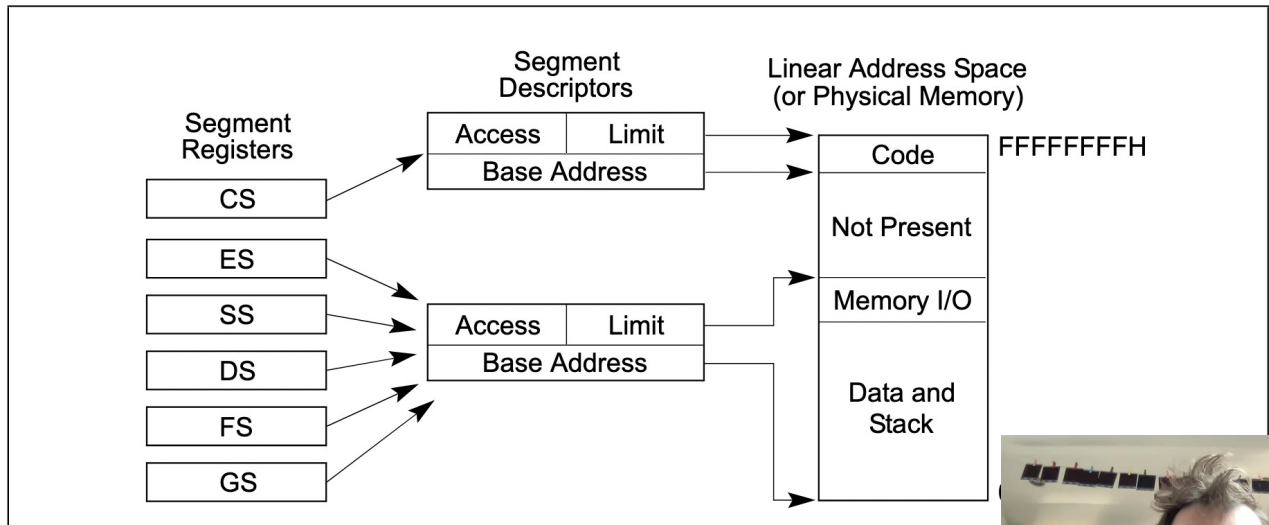


Figure 3-3. Protected Flat Model



Multi-segment model

- Lots of segments for every program and process
- Everything isolated from each other

(This is the one we use)

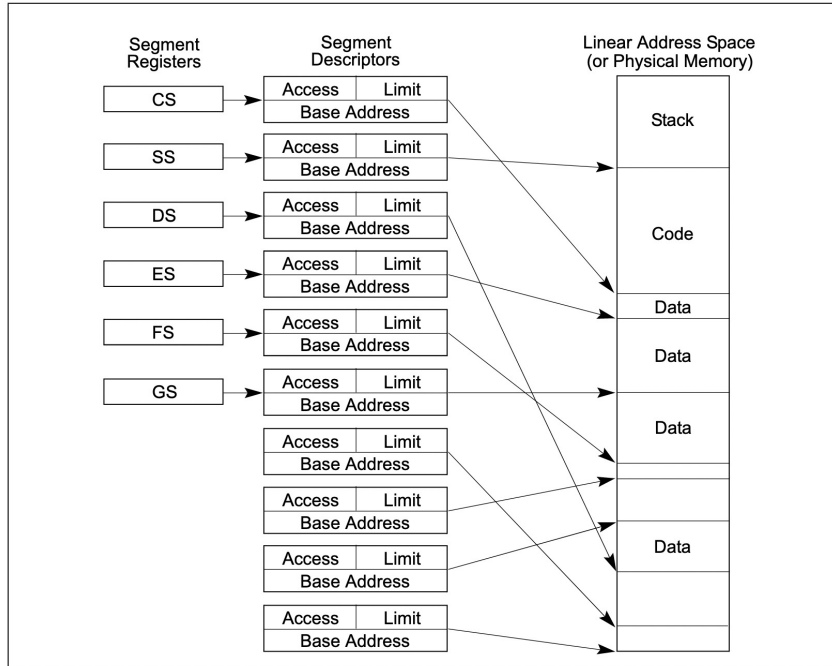


Figure 3-4. Multi-Segment Model



Next problem!

So far segmentation lets us divide up linear memory (the big virtual address space) into smaller areas, per process/ring/whatever

- Limited scope for changing segments within userland
- Limited permissions
 - (read only, execute only, read and execute, read and write... *a few more*)

Sometimes we'd like even finer grained permissions...

- How are we going to do this then?



Paging

Segmentation splits up the virtual address space...

- ...Which is then mapped onto page entries which map to physical address space
- ...And a whole load more tables and permissions

Kernel is allowed to switch pages in and out of memory

- If a process tries to access memory in a different page a PAGEFAULT trap occurs and the kernel can decide what to do
- (*Probably* swap in the right page for the process and let it continue... which is slow)



Paging rules

- Pages are *at least* 4KB on every platform you will ever care about
- Pages cannot span different physical chips
 - The memory has to be physically contiguous
- Pages can configure how processes access physical memory
 - i.e. are writes batched or done directly/rollback
 - (See *Spectre and Meltdown* in COMSM0049 ;-)
- You can disable paging if you *really* want...
 - But then you lose virtual memory and have to deal with the physical address space, and it makes everything worse
 - (This is necessary for some low-level, BIOS/EFI level code)



Paging permissions

- For each page you can set a bunch of extra permission *if your CPU supports it*
 - Almost all architectures have some kind of permissions available
- The big one you need to know about is W^X
 - If you can *write* to memory you shouldn't be able to *execute* it
 - Stops a malicious user finding a buffer overflow, injecting a program, and trivially executing it...ish
 - But it also slows program loading down...



JITing and W^X

- Say I have a program that takes code written by users, and then compiles and runs it on the fly...
- Oh and the code can be dynamically generated and change at short notice
 - Like the JavaScript engine running JS in your web browser
- What has to happen to make this work?



JITing and W^X is slow

1. Get the code chunk you want to compile into memory
2. Compile it and write that code into memory
3. Stick a bit of code on the end to go fetch the next bit of code to compile and repeat
4. Run it

If we have W^X bits we need to make an extra system call between 3 and 4, and before 1 to ensure that the area of memory we're using to load our code into is writable or executable between every chunk.

2 extra syscalls per *Just In Time-compiled* chunk...



We're done!

- Congratulations we've reached the end!
- We're all new lecturers here...
- This was our first year running this unit (and the second year its ever run)...
- Thank you for baring with us!
- Hopefully you enjoyed it...
- Feel free to grab us and chat if you want to talk operating systems and security

bristol.ac.uk



Exam

We've set the exam...

- 8 short questions... 5 marks each
- 4 long questions... 15 marks each
- Equally distributed over the content and reading for the lectures
- Open book

Tips:

- If a question is worth 5 marks give me 5 things to tick
- If a question says *debate*, then make an argument both ways or at least say why the other way is wrong
- Remember that we want you to do well!

bristol.ac.uk

