



Computer Systems B

COMS20012

Introduction to Operating Systems and Security

bristol.ac.uk

Thread Switching

bristol.ac.uk



Previously in COMS20012

- On thread switch
 - Save outgoing thread register values
 - Save outgoing thread stack pointer
 - Load incoming thread stack pointer
 - Load incoming thread register values
- Fairly simple

bristol.ac.uk

Implementing thread switch in OS/161 1/2

```
switchframe_switch:
    addi sp, sp, -40

    sw   ra, 36(sp)
    sw   gp, 32(sp)
    sw   s8, 28(sp)
    sw   s6, 24(sp)
    sw   s5, 20(sp)
    sw   s4, 16(sp)
    sw   s3, 12(sp)
    sw   s2, 8(sp)
    sw   s1, 4(sp)
    sw   s0, 0(sp)

    sw   sp, 0(a0)
```

bristol.ac.uk

Implementing thread switch in OS/161 1/2

switchframe_switch:

```
    addi sp, sp, -40    }    Allocate stack space for 10 registers
```

```
    sw   ra, 36(sp)
    sw   gp, 32(sp)
    sw   s8, 28(sp)
    sw   s6, 24(sp)
    sw   s5, 20(sp)
    sw   s4, 16(sp)
    sw   s3, 12(sp)
    sw   s2, 8(sp)
    sw   s1, 4(sp)
    sw   s0, 0(sp)

    sw   sp, 0(a0)
```

bristol.ac.uk

Implementing thread switch in OS/161 1/2

switchframe_switch:

addi sp, sp, -40

```
sw ra, 36(sp)
sw gp, 32(sp)
sw s8, 28(sp)
sw s6, 24(sp)
sw s5, 20(sp)
sw s4, 16(sp)
sw s3, 12(sp)
sw s2, 8(sp)
sw s1, 4(sp)
sw s0, 0(sp)
```

} Save the registers value to the stack

sw sp, 0(a0)

bristol.ac.uk

Implementing thread switch in OS/161 1/2

switchframe_switch:

addi sp, sp, -40

sw ra, 36(sp)

sw gp, 32(sp)

sw s8, 28(sp)

sw s6, 24(sp)

sw s5, 20(sp)

sw s4, 16(sp)

sw s3, 12(sp)

sw s2, 8(sp)

sw s1, 4(sp)

sw s0, 0(sp)

sw sp, 0(a0) } Store outgoing thread stack pointer

bristol.ac.uk

Implementing thread switch in OS/161 2/2

```
lw sp, 0(a1)
nop
lw s0, 0(sp)
lw s1, 4(sp)
lw s2, 8(sp)
lw s3, 12(sp)
lw s4, 16(sp)
lw s5, 20(sp)
lw s6, 24(sp)
lw s8, 28(sp)
lw gp, 32(sp)
lw ra, 36(sp)
nop
j ra
addi sp, sp, 40
.end switchframe_switch
```

bristol.ac.uk

Implementing thread switch in OS/161 2/2

```
lw sp, 0(at) } Load incoming thread stack pointer
nop
lw s0, 0(sp)
lw s1, 4(sp)
lw s2, 8(sp)
lw s3, 12(sp)
lw s4, 16(sp)
lw s5, 20(sp)
lw s6, 24(sp)
lw s8, 28(sp)
lw gp, 32(sp)
lw ra, 36(sp)
nop
j ra
addi sp, sp, 40
and switchframe_switch
```

bristol.ac.uk

Implementing thread switch in OS/161 2/2

```
lw sp, 0(e1)
nop
lw s0, 0(sp)
lw s1, 4(sp)
lw s2, 8(sp)
lw s3, 12(sp)
lw s4, 16(sp)
lw s5, 20(sp)
lw s6, 24(sp)
lw s8, 28(sp)
lw gp, 32(sp)
lw ra, 36(sp)
nop
j ra
addi sp, sp, 40
and switchframe_switch
```

} Load incoming thread register values

bristol.ac.uk

Implementing thread switch in OS/161 2/2

```
lw sp, 0(e1)
nop
lw s0, 0(sp)
lw s1, 4(sp)
lw s2, 8(sp)
lw s3, 12(sp)
lw s4, 16(sp)
lw s5, 20(sp)
lw s6, 24(sp)
lw s8, 28(sp)
lw gp, 32(sp)
lw ra, 36(sp)
nop
j ra } Return
addi sp, sp, 40
and switchframe_switch
```

bristol.ac.uk

Implementing thread switch in OS/161 2/2

```
lw sp, 0(e1)
nop
lw s0, 0(sp)
lw s1, 4(sp)
lw s2, 8(sp)
lw s3, 12(sp)
lw s4, 16(sp)
lw s5, 20(sp)
lw s6, 24(sp)
lw s8, 28(sp)
lw gp, 32(sp)
lw ra, 36(sp)
nop
j ra
addi sp, sp, 40 } Shrink stack size (mirror our first instruction earlier)
and switchframe_switch
```

bristol.ac.uk

Implementing thread switch in OS/161 2/2

```
lw sp, 0(a1)
nop
lw s0, 0(sp)
lw s1, 4(sp)
lw s2, 8(sp)
lw s3, 12(sp)
lw s4, 16(sp)
lw s5, 20(sp)
lw s6, 24(sp)
lw s8, 28(sp)
lw gp, 32(sp)
lw ra, 36(sp)
nop
j ra
addi sp, sp, 40
.end switchframe_switch
```

bristol.ac.uk

Wait a minute!

- Something happens after the return?!
- ... and what about those “nop”s?!



bristol.ac.uk

MIPS Delay Slot Instruction

- Instruction executed without the effect of the previous instruction
- Present on load, jump, branch etc.
- You can use nop if the effect is necessary to carry the next instruction
- ... or do something useful

bristol.ac.uk

Where to find this code?

- `kern/arch/mips/thread/switch.S`

bristol.ac.uk

Thread switch on x86

```
.globl switch
switch:

    movl 4(%esp), %eax
    movl 8(%esp), %edx

    pushl %ebp
    pushl %ebx
    pushl %esi
    pushl %edi

    movl %esp, %eax
    movl %edx, %esp

    popl %edi
    popl %esi
    popl %ebx
    popl %ebp
    ret
```

bristol.ac.uk

Thread switch on x86

global switch

switch:

```
movl 4(%esp), %eax  
movl 8(%esp), %edx } Expand stack
```

pushl %eax

pushl %edx

pushl %eax

pushl %edx

movl %eax, %eax

movl %edx, %edx

popl %eax

popl %edx

popl %eax

popl %edx

ret

bristol.ac.uk

Thread switch on x86

```
global switch
switch:
    movl 4(%esp), %eax
    movl 8(%esp), %edx

    pushl %ebp
    pushl %ebx
    pushl %esi
    pushl %edi
    } Save registers

    movl %eax, %ecx
    movl %edx, %ecx

    popl %edi
    popl %esi
    popl %ebx
    popl %ebp
    ret
```

bristol.ac.uk

Thread switch on x86

```
global switch
extern
    movl 4(%esp), %eax
    movl 8(%esp), %edx

    pushl %edx
    pushl %eax
    pushl %eax
    pushl %edx

    movl %esp, (%eax)
    movl %edx, %esp
    } Switch stack

    popl %edx
    popl %eax
    popl %edx
    popl %edx
    ret
```

bristol.ac.uk

Thread switch on x86

```
global switch
switch:
    movl 4(%esp), %eax
    movl 8(%esp), %edx

    pushl %edx
    pushl %eax
    pushl %eax
    pushl %eax
    pushl %edx

    movl %eax, %ecx
    movl %edx, %ecx

    popl %edi
    popl %esi
    popl %ebx
    popl %ebp
    ret
```

} Load saved register values

bristol.ac.uk

Thread switch on x86

```
global switch
switch:
    movl $1, %eax
    movl $2, %eax

    pushl %eax
    pushl %eax
    pushl %eax
    pushl %eax

    movl %eax, %eax
    movl %eax, %eax

    popl %eax
    popl %eax
    popl %eax
    popl %eax
    ret } Return
```

bristol.ac.uk

Thread switch on x86

- Different implementation
- High-level logic exactly the same
- Will be the same on all architectures

bristol.ac.uk

What causes context switches?

- Thread switch is a “type” of context switch (see next video)
- Running thread **yield**
 - Voluntarily let another thread run
- Running thread call **thread_exit**
 - Voluntarily stops
- Running thread **blocks**, via a call to **wchan_sleep**
 - We saw that in the semaphore video! (Week 5 Video 6)
- Running thread is **pre-empted**
 - Running thread involuntarily stopped by an **interrupt** (more in the next video)
 - Could be hardware events “**trap**” (Week 5 video 3)
 - ... or the timer used to implement the **scheduler** (more on this next week)

bristol.ac.uk



Thank you

bristol.ac.uk

